

# Transformers 2

Frederick Vandermoeten



# Overview

- Multi-Head Attention
- Complete Transformer Model
- Cross-Attention
- Training
- Transformer Variants
- Optimizations

## Multi-Head Attention

- many Attention Heads in parallel create one Layer

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

- each head uses a linear transformation on Q, K and V

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

## Output Matrix Splitting - Example GPT3

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

$$d_{\text{Embedding}} = d_E = 12,288; d_k = 128$$

$$W_i^Q \text{ and } W_i^K \in \mathbb{R}^{d_E \times d_k}$$

$$d_e * d_k = 12,288 * 128 = 1,572,864$$

$$W_{i,pre}^V \in \mathbb{R}^{d_E \times d_E}$$

$$d_e * d_e = 12,288 * 12,288 = 150,994,944$$

## Output Matrix Splitting - Example GPT3

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

$$W_i^Q \text{ and } W_i^K \in \mathbb{R}^{d_E \times d_k}$$

$$d_e * d_k = 12,288 * 128 = 1,572,864$$

$$W_{i,pre}^V \in \mathbb{R}^{d_E \times d_E}$$


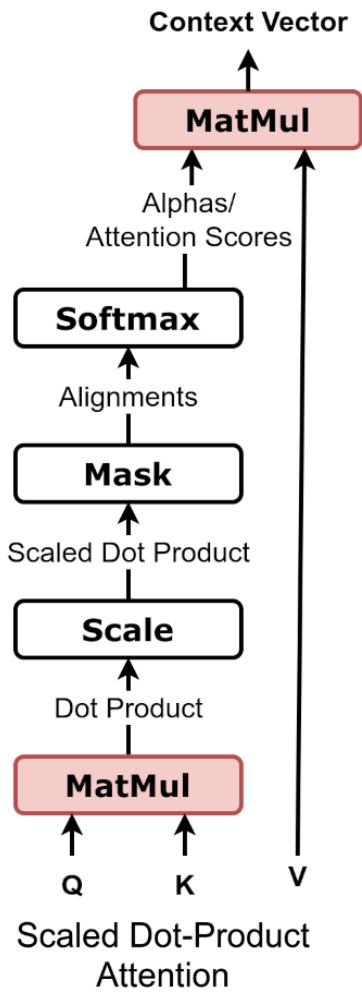
$$d_e * d_e = 12,288 * 12,288 = 150,994,944$$

We split  $W_{i,pre}^V$  into  $W_i^V \in \mathbb{R}^{d_E \times d_v}$  and  $W_i^O \in \mathbb{R}^{d_v \times d_E}$  with  $W_{i,pre}^V = W_i^V W_i^O$

If we use  $d_v = d_k$ , we get the same amount of parameters for  $W_i^V$  and  $W_i^O$  as for  $W_i^Q$  and  $W_i^K$

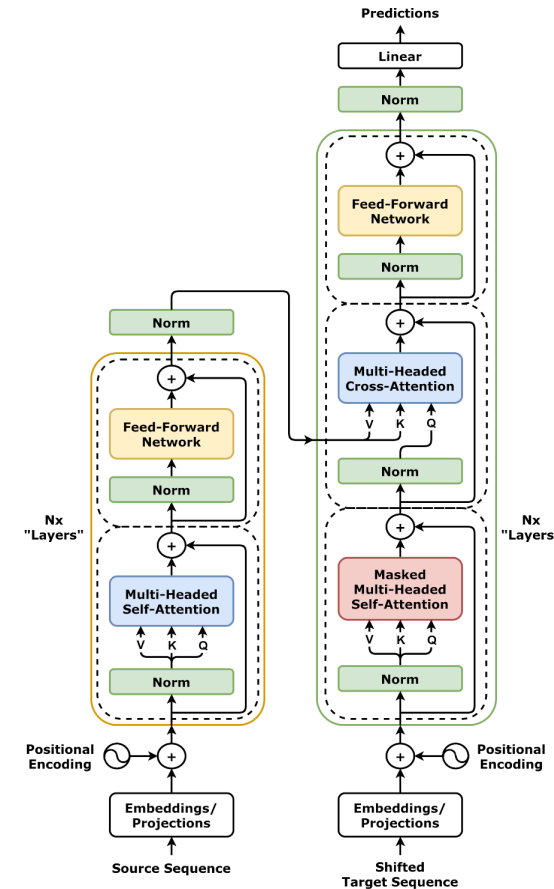
$$W^O = \text{Vertcat}(W_1^O, \dots, W_h^O)$$

# Multi-Head Attention 2




# Complete Transformer Model

- Encoder-Decoder architecture
- Embedding
- Positional Encoding
- Multi-Head Attention
- Feed Forward Network
- Skip Connections
- Layer Normalization
- Cross-Attention



# Cross-Attention

- Enables interaction between encoder and decoder
- Decoder uses encoder output as input
- Query comes from the decoder, key and value from the encoder
- Encoder and decoder could be from different modals (text, image, video)




RAS2024 (modified)



# Training

## Parallelization

- Unlike RNNs, no sequential dependencies between inputs
- Just a bunch of matrix multiplications that can be parallelized very well with modern hardware (GPUs, TPUs)
- We still get order information due to the positional encoding
- Efficient use of training data



## Training: Example GPT-2 vs GPT-3

### GPT-2

- 1.5B parameters
- 40 Gb of data

### GPT-3

- 176B parameters
- 570 Gb of data

## Transformer Variants

Transformers were initially designed for NLP, but have since been applied to other tasks as well.

# Text Transformers

## GPT - Generative Pre-trained Transformer

- Decoder-only Transformer
- Used for text generation, summarization, question answering, etc.

## BERT - Bidirectional Encoder Representations from Transformers

- Encoder-only Transformer
- Used for text classification, sentiment analysis, etc.

## T5 - Text-to-Text Transfer Transformer

- Encoder-Decoder Transformer
- Less specialized so it needs fine-tuning, useful for e.g. translation

# Vision Transformers

## ViT - Vision Transformer

- Images are encoded in patches, then the patches are fed into the Transformer
- Self-Attention is used to learn the global relationships
- Used for image classification, object detection, etc.

## Stable Diffusion

- Diffusion models are encoder-decoder models that use noise for the hidden state
- Cross-Attention is used to integrate text prompts
- Used for image generation, image inpainting, etc.

# Audio Transformers

## Wav2Vec

- Operates directly on waveforms
- Speech recognition, especially for low-resource languages

## Audio Spectrogram Transformer

- Uses ViT on spectrograms
- Audio classification (genre, instrument, etc.), speech emotion recognition, etc.

# Optimizations

## Sparse Transformers

- Attention Layers grow  $O(n^2)$  for the context size  $n$
- most results of  $\text{softmax}(\frac{Q^T K}{\sqrt{d_k}})$  are close to zero
- different ways to reduce the scaling by only considering close values or random feature attention

## Model Distillation

- Train a smaller, faster model with data generated by a larger, slower model
- Teacher/ Student Model

## Conclusion

- Transformers revolutionized NLP
- Attention is what drives their success
- They are versatile and can be used for many different tasks
- They scale well to large datasets
- Work in progress



## Image Sources

All not locally credited images are from [dvgodoy: Deep Learning Visuals CC BY](#)

## Main Conceptual Sources

- [Vaswani, A. "Attention is all you need." Advances in Neural Information Processing Systems \(2017\).](#)
- [Raschka: Understanding and coding self-attention](#)
- [3Blue1Brown: Neural Networks Playlist](#)
- [SBert Wiki](#)